

Week 5 - Neural Networks - Classification - Python

October 10, 2018

1 Data Warehousing and Data Mining

1.1 Labs

1.1.1 Prepared by Gilroy Gordon

Contact Information SCIT ext. 3643

ggordon@utech@gmail.com

gilroy.gordon@utech.edu.jm

1.1.2 Week 5 - Neural Networks in Python

Additional Reference Resources:

http://scikit-learn.org/stable/modules/neural_networks_supervised.html

<https://machinelearningmastery.com/how-to-configure-the-number-of-layers-and-nodes-in-a-neural-network/>

1.2 Objectives

- > Data Preprocessing
 - > Min Max Scaling
- > Data Transformation
- > Data Mining
 - > Neural Networks (Classification and Estimation)
- > Model Evaluation and Prediction
 - > Train/Test Split - 70/30
- > Presentation
 - > Plots

1.3 Import required libraries and acquire data

```
In [38]: # import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [39]: data_path = './data/hr_data.csv' # Path to data file
data = pd.read_csv(data_path)
data.head(15)
```

```
Out[39]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	\
0	0.38	0.53	2	157	
1	0.80	0.86	5	262	
2	0.11	0.88	7	272	
3	0.72	0.87	5	223	
4	0.37	0.52	2	159	
5	0.41	0.50	2	153	
6	0.10	0.77	6	247	
7	0.92	0.85	5	259	
8	0.89	1.00	5	224	
9	0.42	0.53	2	142	
10	0.45	0.54	2	135	
11	0.11	0.81	6	305	
12	0.84	0.92	4	234	
13	0.41	0.55	2	148	
14	0.36	0.56	2	137	

	time_spend_company	Work_accident	left	promotion_last_5years	sales	\
0	3	0	1	0	sales	
1	6	0	1	0	sales	
2	4	0	1	0	sales	
3	5	0	1	0	sales	
4	3	0	1	0	sales	
5	3	0	1	0	sales	
6	4	0	1	0	sales	
7	5	0	1	0	sales	
8	5	0	1	0	sales	
9	3	0	1	0	sales	
10	3	0	1	0	sales	
11	4	0	1	0	sales	
12	5	0	1	0	sales	
13	3	0	1	0	sales	
14	3	0	1	0	sales	

	salary
0	low
1	medium
2	medium
3	low
4	low
5	low
6	low
7	low
8	low

```

9      low
10     low
11     low
12     low
13     low
14     low

```

```
In [40]: # What columns are in the data set ? Do they have spaces that I should consider
data.columns
```

```
Out[40]: Index(['satisfaction_level', 'last_evaluation', 'number_project',
               'average_monthly_hours', 'time_spend_company', 'Work_accident', 'left',
               'promotion_last_5years', 'sales', 'salary'],
              dtype='object')
```

1.4 Aim: Can we determine a person's Department based on the other factors?

department (column currently named sales) = a(last_evaluation) + b(number_project) + c(average_monthly_hours) + d(time_spend_company)

The coefficients a-d, what are they? What is the relationship between the variables? Does multicollinearity exist?

I have created a function below create_label_encoder_dict to assist with this. The function accepts a dataframe object and uses the LabelEncoder class from sklearn.preprocessing to encode (dummy encoding) or transform non-numerical columns to numbers. Finally it returns a dictionary object of all the encoders created for each column.

The LabelEncoder is a useful resource as it not only automatically transforms all values in a column but also keeps a track of what values were transformed from. i.e. It will change all Female to 0 and all Male to 1

```
In [41]: def create_label_encoder_dict(df):
          from sklearn.preprocessing import LabelEncoder

          label_encoder_dict = {}
          for column in df.columns:
              # Only create encoder for categorical data types
              if not np.issubdtype(df[column].dtype, np.number) and column != 'Age':
                  label_encoder_dict[column] = LabelEncoder().fit(df[column])
          return label_encoder_dict

In [42]: label_encoders = create_label_encoder_dict(data)
          print("Encoded Values for each Label")
          print("="*32)
          for column in label_encoders:
              print("="*32)
              print('Encoder(%s) = %s' % (column, label_encoders[column].classes_))
              print(pd.DataFrame([range(0, len(label_encoders[column].classes_))], columns=label_e
```

```
Encoded Values for each Label
=====
```

```

=====
Encoder(sales) = ['IT' 'RandD' 'accounting' 'hr' 'management' 'marketing' 'product_mng'
                  'sales' 'support' 'technical']
                  Encoded Values
IT                0
RandD             1
accounting        2
hr               3
management        4
marketing         5
product_mng       6
sales             7
support           8
technical         9
=====
Encoder(salary) = ['high' 'low' 'medium']
                  Encoded Values
high              0
low               1
medium           2

```

```

In [43]: # Apply each encoder to the data set to obtain transformed values
data2 = data.copy() # create copy of initial data set
for column in data2.columns:
    if column in label_encoders:
        data2[column] = label_encoders[column].transform(data2[column])

print("Transformed data set")
print("="*32)
data2.head(15)

```

Transformed data set

```
=====
```

```

Out[43]:

```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	\
0	0.38	0.53	2	157	
1	0.80	0.86	5	262	
2	0.11	0.88	7	272	
3	0.72	0.87	5	223	
4	0.37	0.52	2	159	
5	0.41	0.50	2	153	
6	0.10	0.77	6	247	
7	0.92	0.85	5	259	
8	0.89	1.00	5	224	
9	0.42	0.53	2	142	
10	0.45	0.54	2	135	

11	0.11	0.81	6	305
12	0.84	0.92	4	234
13	0.41	0.55	2	148
14	0.36	0.56	2	137

	time_spend_company	Work_accident	left	promotion_last_5years	sales	\
0	3	0	1	0	7	
1	6	0	1	0	7	
2	4	0	1	0	7	
3	5	0	1	0	7	
4	3	0	1	0	7	
5	3	0	1	0	7	
6	4	0	1	0	7	
7	5	0	1	0	7	
8	5	0	1	0	7	
9	3	0	1	0	7	
10	3	0	1	0	7	
11	4	0	1	0	7	
12	5	0	1	0	7	
13	3	0	1	0	7	
14	3	0	1	0	7	

	salary
0	1
1	2
2	2
3	1
4	1
5	1
6	1
7	1
8	1
9	1
10	1
11	1
12	1
13	1
14	1

```
In [44]: # separate our data into dependent (Y) and independent(X) variables
X_data = data2[['last_evaluation', 'number_project', 'average_monthly_hours', 'time_spend_c
Y_data = data2['sales'] # actually department column
```

1.5 70/30 Train Test Split

We will split the data using a 70/30 split. i.e. 70% of the data will be randomly chosen to train the model and 30% will be used to evaluate the model

```

In [45]: from sklearn.model_selection import train_test_split
         X_train, X_test, y_train, y_test = train_test_split(X_data, Y_data, test_size=0.30)

In [46]: from sklearn.neural_network import MLPClassifier

In [106]: # Create an instance of linear regression
          reg = MLPClassifier()
          #reg = MLPClassifier(hidden_layer_sizes=(8,120))

In [107]: reg.fit(X_train,y_train)

Out[107]: MLPClassifier(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,
                        beta_2=0.999, early_stopping=False, epsilon=1e-08,
                        hidden_layer_sizes=(100,), learning_rate='constant',
                        learning_rate_init=0.001, max_iter=200, momentum=0.9,
                        nesterovs_momentum=True, power_t=0.5, random_state=None,
                        shuffle=True, solver='adam', tol=0.0001, validation_fraction=0.1,
                        verbose=False, warm_start=False)

In [76]: help(MLPClassifier)

```

Help on class MLPClassifier in module sklearn.neural_network.multilayer_perceptron:

```

class MLPClassifier(BaseMultilayerPerceptron, sklearn.base.ClassifierMixin)
|   Multi-layer Perceptron classifier.
|
|   This model optimizes the log-loss function using LBFGS or stochastic
|   gradient descent.
|
|   .. versionadded:: 0.18
|
|   Parameters
|   -----
|   hidden_layer_sizes : tuple, length = n_layers - 2, default (100,)
|       The ith element represents the number of neurons in the ith
|       hidden layer.
|
|   activation : {'identity', 'logistic', 'tanh', 'relu'}, default 'relu'
|       Activation function for the hidden layer.
|
|       - 'identity', no-op activation, useful to implement linear bottleneck,
|         returns  $f(x) = x$ 
|
|       - 'logistic', the logistic sigmoid function,
|         returns  $f(x) = 1 / (1 + \exp(-x))$ .
|
|       - 'tanh', the hyperbolic tan function,
|         returns  $f(x) = \tanh(x)$ .
|

```

```

|     - 'relu', the rectified linear unit function,
|       returns  $f(x) = \max(0, x)$ 
|
| solver : {'lbfgs', 'sgd', 'adam'}, default 'adam'
|   The solver for weight optimization.
|
|     - 'lbfgs' is an optimizer in the family of quasi-Newton methods.
|
|     - 'sgd' refers to stochastic gradient descent.
|
|     - 'adam' refers to a stochastic gradient-based optimizer proposed
|       by Kingma, Diederik, and Jimmy Ba
|
| Note: The default solver 'adam' works pretty well on relatively
| large datasets (with thousands of training samples or more) in terms of
| both training time and validation score.
| For small datasets, however, 'lbfgs' can converge faster and perform
| better.
|
| alpha : float, optional, default 0.0001
|   L2 penalty (regularization term) parameter.
|
| batch_size : int, optional, default 'auto'
|   Size of minibatches for stochastic optimizers.
|   If the solver is 'lbfgs', the classifier will not use minibatch.
|   When set to "auto", `batch_size=min(200, n_samples)`
|
| learning_rate : {'constant', 'invscaling', 'adaptive'}, default 'constant'
|   Learning rate schedule for weight updates.
|
|     - 'constant' is a constant learning rate given by
|       'learning_rate_init'.
|
|     - 'invscaling' gradually decreases the learning rate ``learning_rate``
|       at each time step 't' using an inverse scaling exponent of 'power_t'.
|        $\text{effective\_learning\_rate} = \text{learning\_rate\_init} / \text{pow}(t, \text{power\_t})$ 
|
|     - 'adaptive' keeps the learning rate constant to
|       'learning_rate_init' as long as training loss keeps decreasing.
|       Each time two consecutive epochs fail to decrease training loss by at
|       least tol, or fail to increase validation score by at least tol if
|       'early_stopping' is on, the current learning rate is divided by 5.
|
|   Only used when ``solver='sgd'``.
|
| learning_rate_init : double, optional, default 0.001
|   The initial learning rate used. It controls the step-size
|   in updating the weights. Only used when solver='sgd' or 'adam'.

```

`power_t` : double, optional, default 0.5
 The exponent for inverse scaling learning rate.
 It is used in updating effective learning rate when the `learning_rate` is set to 'invscaling'. Only used when `solver='sgd'`.

`max_iter` : int, optional, default 200
 Maximum number of iterations. The solver iterates until convergence (determined by 'tol') or this number of iterations. For stochastic solvers ('sgd', 'adam'), note that this determines the number of epochs (how many times each data point will be used), not the number of gradient steps.

`shuffle` : bool, optional, default True
 Whether to shuffle samples in each iteration. Only used when `solver='sgd'` or 'adam'.

`random_state` : int, RandomState instance or None, optional, default None
 If int, `random_state` is the seed used by the random number generator;
 If RandomState instance, `random_state` is the random number generator;
 If None, the random number generator is the RandomState instance used by ``np.random``.

`tol` : float, optional, default 1e-4
 Tolerance for the optimization. When the loss or score is not improving by at least `tol` for two consecutive iterations, unless ``learning_rate`` is set to 'adaptive', convergence is considered to be reached and training stops.

`verbose` : bool, optional, default False
 Whether to print progress messages to stdout.

`warm_start` : bool, optional, default False
 When set to True, reuse the solution of the previous call to fit as initialization, otherwise, just erase the previous solution.

`momentum` : float, default 0.9
 Momentum for gradient descent update. Should be between 0 and 1. Only used when `solver='sgd'`.

`nesterovs_momentum` : boolean, default True
 Whether to use Nesterov's momentum. Only used when `solver='sgd'` and `momentum > 0`.

`early_stopping` : bool, default False
 Whether to use early stopping to terminate training when validation score is not improving. If set to true, it will automatically set


```

|         aside 10% of training data as validation and terminate training when
|         validation score is not improving by at least tol for two consecutive
|         epochs.
|         Only effective when solver='sgd' or 'adam'
|
| validation_fraction : float, optional, default 0.1
|         The proportion of training data to set aside as validation set for
|         early stopping. Must be between 0 and 1.
|         Only used if early_stopping is True
|
| beta_1 : float, optional, default 0.9
|         Exponential decay rate for estimates of first moment vector in adam,
|         should be in [0, 1). Only used when solver='adam'
|
| beta_2 : float, optional, default 0.999
|         Exponential decay rate for estimates of second moment vector in adam,
|         should be in [0, 1). Only used when solver='adam'
|
| epsilon : float, optional, default 1e-8
|         Value for numerical stability in adam. Only used when solver='adam'
|
| Attributes
| -----
| classes_ : array or list of array of shape (n_classes,)
|         Class labels for each output.
|
| loss_ : float
|         The current loss computed with the loss function.
|
| coefs_ : list, length n_layers - 1
|         The ith element in the list represents the weight matrix corresponding
|         to layer i.
|
| intercepts_ : list, length n_layers - 1
|         The ith element in the list represents the bias vector corresponding to
|         layer i + 1.
|
| n_iter_ : int,
|         The number of iterations the solver has ran.
|
| n_layers_ : int
|         Number of layers.
|
| n_outputs_ : int
|         Number of outputs.
|
| out_activation_ : string
|         Name of the output activation function.

```

```

|
| Notes
| -----
| MLPClassifier trains iteratively since at each time step
| the partial derivatives of the loss function with respect to the model
| parameters are computed to update the parameters.
|
| It can also have a regularization term added to the loss function
| that shrinks model parameters to prevent overfitting.
|
| This implementation works with data represented as dense numpy arrays or
| sparse scipy arrays of floating point values.
|
| References
| -----
| Hinton, Geoffrey E.
|     "Connectionist learning procedures." Artificial intelligence 40.1
|     (1989): 185-234.
|
| Glorot, Xavier, and Yoshua Bengio. "Understanding the difficulty of
|     training deep feedforward neural networks." International Conference
|     on Artificial Intelligence and Statistics. 2010.
|
| He, Kaiming, et al. "Delving deep into rectifiers: Surpassing human-level
|     performance on imagenet classification." arXiv preprint
|     arXiv:1502.01852 (2015).
|
| Kingma, Diederik, and Jimmy Ba. "Adam: A method for stochastic
|     optimization." arXiv preprint arXiv:1412.6980 (2014).
|
| Method resolution order:
|     MLPClassifier
|     BaseMultilayerPerceptron
|     abc.NewBase
|     sklearn.base.BaseEstimator
|     sklearn.base.ClassifierMixin
|     builtins.object
|
| Methods defined here:
|
|     __init__(self, hidden_layer_sizes=(100,), activation='relu', solver='adam', alpha=0.0001, ba
|         Initialize self.  See help(type(self)) for accurate signature.
|
|     fit(self, X, y)
|         Fit the model to data matrix X and target(s) y.
|
|     Parameters
|     -----

```

```

|     X : array-like or sparse matrix, shape (n_samples, n_features)
|         The input data.
|
|     y : array-like, shape (n_samples,) or (n_samples, n_outputs)
|         The target values (class labels in classification, real numbers in
|         regression).
|
|     Returns
|     -----
|     self : returns a trained MLP model.
|
| predict(self, X)
|     Predict using the multi-layer perceptron classifier
|
|     Parameters
|     -----
|     X : {array-like, sparse matrix}, shape (n_samples, n_features)
|         The input data.
|
|     Returns
|     -----
|     y : array-like, shape (n_samples,) or (n_samples, n_classes)
|         The predicted classes.
|
| predict_log_proba(self, X)
|     Return the log of probability estimates.
|
|     Parameters
|     -----
|     X : array-like, shape (n_samples, n_features)
|         The input data.
|
|     Returns
|     -----
|     log_y_prob : array-like, shape (n_samples, n_classes)
|         The predicted log-probability of the sample for each class
|         in the model, where classes are ordered as they are in
|         `self.classes_`. Equivalent to log(predict_proba(X))
|
| predict_proba(self, X)
|     Probability estimates.
|
|     Parameters
|     -----
|     X : {array-like, sparse matrix}, shape (n_samples, n_features)
|         The input data.
|
|     Returns

```

```

|         -----
|         y_prob : array-like, shape (n_samples, n_classes)
|             The predicted probability of the sample for each class in the
|             model, where classes are ordered as they are in `self.classes_`.
|
|         -----
|         Data descriptors defined here:
|
|         partial_fit
|             Fit the model to data matrix X and target y.
|
|             Parameters
|             -----
|             X : {array-like, sparse matrix}, shape (n_samples, n_features)
|                 The input data.
|
|             y : array-like, shape (n_samples,)
|                 The target values.
|
|             classes : array, shape (n_classes)
|                 Classes across all calls to partial_fit.
|                 Can be obtained via `np.unique(y_all)`, where y_all is the
|                 target vector of the entire dataset.
|                 This argument is required for the first call to partial_fit
|                 and can be omitted in the subsequent calls.
|                 Note that y doesn't need to contain all labels in `classes`.
|
|             Returns
|             -----
|             self : returns a trained MLP model.
|
|         -----
|         Data and other attributes defined here:
|
|         __abstractmethods__ = frozenset()
|
|         -----
|         Methods inherited from sklearn.base.BaseEstimator:
|
|         __getstate__(self)
|
|         __repr__(self)
|             Return repr(self).
|
|         __setstate__(self, state)
|
|         get_params(self, deep=True)
|             Get parameters for this estimator.

```

```

Parameters
-----
deep : boolean, optional
    If True, will return the parameters for this estimator and
    contained subobjects that are estimators.

Returns
-----
params : mapping of string to any
    Parameter names mapped to their values.

set_params(self, **params)
    Set the parameters of this estimator.

    The method works on simple estimators as well as on nested objects
    (such as pipelines). The latter have parameters of the form
    ``<component>__<parameter>`` so that it's possible to update each
    component of a nested object.

Returns
-----
self

-----
Data descriptors inherited from sklearn.base.BaseEstimator:

__dict__
    dictionary for instance variables (if defined)

__weakref__
    list of weak references to the object (if defined)

-----
Methods inherited from sklearn.base.ClassifierMixin:

score(self, X, y, sample_weight=None)
    Returns the mean accuracy on the given test data and labels.

    In multi-label classification, this is the subset accuracy
    which is a harsh metric since you require for each sample that
    each label set be correctly predicted.

Parameters
-----
X : array-like, shape = (n_samples, n_features)
    Test samples.

```

```

|     y : array-like, shape = (n_samples) or (n_samples, n_outputs)
|         True labels for X.
|
|     sample_weight : array-like, shape = [n_samples], optional
|         Sample weights.
|
| Returns
| -----
| score : float
|     Mean accuracy of self.predict(X) wrt. y.

```

```
In [108]: reg.n_layers_ # Number of layers utilized
```

```
Out[108]: 3
```

```
In [109]: # Make predictions using the testing set
          test_predicted = reg.predict(X_test)
          test_predicted
```

```
Out[109]: array([7, 7, 7, ..., 7, 7, 7])
```

```
In [111]: data3 = X_test.copy()
          data3['predicted_department']=test_predicted
          data3['predicted_department_en']=label_encoders['sales'].inverse_transform(test_predicted)
          data3['department']=data3['sales']=y_test
          data3['department_en']=label_encoders['sales'].inverse_transform(y_test)
          data3.head()
```

```

/usr/local/lib/python3.5/dist-packages/sklearn/preprocessing/label.py:151: DeprecationWarning: T
if diff:
/usr/local/lib/python3.5/dist-packages/sklearn/preprocessing/label.py:151: DeprecationWarning: T
if diff:

```

```
Out[111]:
```

	last_evaluation	number_project	average_monthly_hours	\
7130	0.82	3	134	
13591	0.62	6	225	
14003	0.74	3	265	
3710	0.46	2	169	
12310	0.93	7	305	

	time_spend_company	predicted_department	predicted_department_en	\
7130	3	7	sales	
13591	6	7	sales	
14003	3	7	sales	
3710	2	7	sales	
12310	4	7	sales	

	department	sales	department_en
7130	2	2	accounting
13591	7	7	sales
14003	8	8	support
3710	8	8	support
12310	9	9	technical

1.6 Evaluation

1.6.1 Building a Confusion Matrix

NB. Data should be split in training and test data. The model built should be evaluated using unseen or test data

```
In [112]: k=(reg.predict(X_test) == y_test) # Determine how many were predicted correctly
```

```
In [113]: k.value_counts()
```

```
Out[113]: False    3232
          True     1268
          Name: sales, dtype: int64
```

```
In [114]: from sklearn.metrics import confusion_matrix
```

```
In [116]: cm=confusion_matrix(y_test, reg.predict(X_test), labels=y_test.unique())
          cm
```

```
Out[116]: array([[ 0, 230,  2,  0,  0,  0,  0,  0,  0,  0],
                 [ 0, 1266,  1,  0,  0,  0,  0,  0,  0,  0],
                 [ 0,  664,  2,  0,  0,  0,  0,  0,  0,  0],
                 [ 0,  829,  1,  0,  0,  0,  0,  0,  0,  0],
                 [ 0,  370,  1,  0,  0,  0,  0,  0,  0,  0],
                 [ 0,  240,  0,  0,  0,  0,  0,  0,  0,  0],
                 [ 0,  256,  0,  0,  0,  0,  0,  0,  0,  0],
                 [ 0,  221,  0,  0,  0,  0,  0,  0,  0,  0],
                 [ 0,  234,  0,  0,  0,  0,  0,  0,  0,  0],
                 [ 0,  182,  1,  0,  0,  0,  0,  0,  0,  0]])
```

```
In [117]: def plot_confusion_matrix(cm, classes,
                                   normalize=False,
                                   title='Confusion matrix',
                                   cmap=plt.cm.Blues):

    import itertools
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if normalize:
```

```

        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

```

```

In [122]: plt.figure(figsize=(9,16))
          plot_confusion_matrix(cm,data['sales'].unique())


```

Confusion matrix, without normalization

```

[[ 0 230  2  0  0  0  0  0  0  0]
 [ 0 1266 1  0  0  0  0  0  0  0]
 [ 0 664  2  0  0  0  0  0  0  0]
 [ 0 829  1  0  0  0  0  0  0  0]
 [ 0 370  1  0  0  0  0  0  0  0]
 [ 0 240  0  0  0  0  0  0  0  0]
 [ 0 256  0  0  0  0  0  0  0  0]
 [ 0 221  0  0  0  0  0  0  0  0]
 [ 0 234  0  0  0  0  0  0  0  0]
 [ 0 182  1  0  0  0  0  0  0  0]]

```

2 With Preprocessing

```
In [123]: def create_min_max_scaler_dict(df):
            from sklearn.preprocessing import MinMaxScaler
            min_max_scaler_dict = {}
            for column in df.columns:
                # Only create encoder for categorical data types
                if np.issubdtype(df[column].dtype, np.number):
                    min_max_scaler_dict[column] = MinMaxScaler().fit(pd.DataFrame(df[column]))
            return min_max_scaler_dict
```

```
In [124]: min_max_scalers = create_min_max_scaler_dict(data)
            print("Min Max Values for each Label")
            print("="*32)
            min_max_scalers
```

```
Min Max Values for each Label
=====
```

```
Out[124]: {'Work_accident': MinMaxScaler(copy=True, feature_range=(0, 1)),
            'average_monthly_hours': MinMaxScaler(copy=True, feature_range=(0, 1)),
            'last_evaluation': MinMaxScaler(copy=True, feature_range=(0, 1)),
            'left': MinMaxScaler(copy=True, feature_range=(0, 1)),
            'number_project': MinMaxScaler(copy=True, feature_range=(0, 1)),
            'promotion_last_5years': MinMaxScaler(copy=True, feature_range=(0, 1)),
            'satisfaction_level': MinMaxScaler(copy=True, feature_range=(0, 1)),
            'time_spend_company': MinMaxScaler(copy=True, feature_range=(0, 1))}
```

```
In [125]: #retrieving a scaler
            time_spend_company_scaler=min_max_scalers['time_spend_company']
```

```
In [126]: time_spend_company_scaler
```

```
Out[126]: MinMaxScaler(copy=True, feature_range=(0, 1))
```

```
In [127]: time_spend_company_scaler.data_max_ #Maximum value
```

```
Out[127]: array([10.])
```

```
In [128]: time_spend_company_scaler.data_min_ # Minimum value
```

```
Out[128]: array([2.])
```

```
In [129]: time_spend_company_scaler.data_range_ # Range = Max- Min
```

```
Out[129]: array([8.])
```

```
In [130]: pd.DataFrame([
    {
        'column':col,
        'min':min_max_scalers[col].data_min_[0],
        'max':min_max_scalers[col].data_max_[0],
        'range':min_max_scalers[col].data_range_[0] } for col in min_max_scalers])
```

```
Out[130]:
```

	column	max	min	range
0	time_spend_company	10	2.00	8.00
1	promotion_last_5years	1	0.00	1.00
2	satisfaction_level	1	0.09	0.91
3	left	1	0.00	1.00
4	average_monthly_hours	310	96.00	214.00
5	number_project	7	2.00	5.00
6	Work_accident	1	0.00	1.00
7	last_evaluation	1	0.36	0.64

```
In [131]: # Apply each scaler to the data set to obtain transformed values
data3 = data2.copy() # create copy of initial data set
for column in data3.columns:
    if column in min_max_scalers:
        data3[column] = min_max_scalers[column].transform(pd.DataFrame(data3[column]))

print("Transformed data set")
print("="*32)
data3.head(15)
```

Transformed data set

=====

```
Out[131]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	\
0	0.318681	0.265625	0.0	0.285047	
1	0.780220	0.781250	0.6	0.775701	
2	0.021978	0.812500	1.0	0.822430	
3	0.692308	0.796875	0.6	0.593458	
4	0.307692	0.250000	0.0	0.294393	
5	0.351648	0.218750	0.0	0.266355	
6	0.010989	0.640625	0.8	0.705607	
7	0.912088	0.765625	0.6	0.761682	
8	0.879121	1.000000	0.6	0.598131	
9	0.362637	0.265625	0.0	0.214953	
10	0.395604	0.281250	0.0	0.182243	
11	0.021978	0.703125	0.8	0.976636	
12	0.824176	0.875000	0.4	0.644860	
13	0.351648	0.296875	0.0	0.242991	
14	0.296703	0.312500	0.0	0.191589	

time_spend_company	Work_accident	left	promotion_last_5years	sales	\
--------------------	---------------	------	-----------------------	-------	---

0	0.125	0	1	0	7
1	0.500	0	1	0	7
2	0.250	0	1	0	7
3	0.375	0	1	0	7
4	0.125	0	1	0	7
5	0.125	0	1	0	7
6	0.250	0	1	0	7
7	0.375	0	1	0	7
8	0.375	0	1	0	7
9	0.125	0	1	0	7
10	0.125	0	1	0	7
11	0.250	0	1	0	7
12	0.375	0	1	0	7
13	0.125	0	1	0	7
14	0.125	0	1	0	7

	salary
0	1
1	2
2	2
3	1
4	1
5	1
6	1
7	1
8	1
9	1
10	1
11	1
12	1
13	1
14	1

In [132]: data3.describe()

```
Out[132]:
```

	satisfaction_level	last_evaluation	number_project	\
count	14999.000000	14999.000000	14999.000000	
mean	0.574542	0.556409	0.360611	
std	0.273220	0.267452	0.246518	
min	0.000000	0.000000	0.000000	
25%	0.384615	0.312500	0.200000	
50%	0.604396	0.562500	0.400000	
75%	0.802198	0.796875	0.600000	
max	1.000000	1.000000	1.000000	

	average_monthly_hours	time_spend_company	Work_accident	left	\
count	14999.000000	14999.000000	14999.000000	14999.000000	
mean	0.490889	0.187279	0.144610	0.238083	

std	0.233379	0.182517	0.351719	0.425924
min	0.000000	0.000000	0.000000	0.000000
25%	0.280374	0.125000	0.000000	0.000000
50%	0.485981	0.125000	0.000000	0.000000
75%	0.696262	0.250000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000

	promotion_last_5years	sales	salary
count	14999.000000	14999.000000	14999.000000
mean	0.021268	5.870525	1.347290
std	0.144281	2.868786	0.625819
min	0.000000	0.000000	0.000000
25%	0.000000	4.000000	1.000000
50%	0.000000	7.000000	1.000000
75%	0.000000	8.000000	2.000000
max	1.000000	9.000000	2.000000

```
In [133]: # separate our data into dependent (Y) and independent(X) variables
X2_data = data3[['last_evaluation', 'number_project', 'average_monthly_hours', 'time_spent_per_project']]
Y2_data = data3['sales']
```

2.1 70/30 Train Test Split

We will split the data using a 70/30 split. i.e. 70% of the data will be randomly chosen to train the model and 30% will be used to evaluate the model

```
In [134]: from sklearn.model_selection import train_test_split
X2_train, X2_test, y2_train, y2_test = train_test_split(X2_data, Y2_data, test_size=0.3)
```

```
In [135]: # Create an instance of linear regression
reg2 = MLPClassifier()
reg2.fit(X2_train, y2_train)
```

```
Out[135]: MLPClassifier(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,
                        beta_2=0.999, early_stopping=False, epsilon=1e-08,
                        hidden_layer_sizes=(100,), learning_rate='constant',
                        learning_rate_init=0.001, max_iter=200, momentum=0.9,
                        nesterovs_momentum=True, power_t=0.5, random_state=None,
                        shuffle=True, solver='adam', tol=0.0001, validation_fraction=0.1,
                        verbose=False, warm_start=False)
```

```
In [136]: reg2.n_layers_
```

```
Out[136]: 3
```

```
In [137]: # Make predictions using the testing set
test2_predicted = reg2.predict(X2_test)
test2_predicted
```

```
Out[137]: array([7, 7, 7, ..., 7, 7, 7])
```

```
In [138]: data4 = X2_test.copy()
          data4['predicted_department']=test2_predicted
          data4['predicted_department_en']=label_encoders['sales'].inverse_transform(test2_predi
          data4['department']=data3['sales']=y2_test
          data4['department_en']=label_encoders['sales'].inverse_transform(y2_test)
          data4.head()
```

```
/usr/local/lib/python3.5/dist-packages/sklearn/preprocessing/label.py:151: DeprecationWarning: T
if diff:
/usr/local/lib/python3.5/dist-packages/sklearn/preprocessing/label.py:151: DeprecationWarning: T
if diff:
```

```
Out[138]:
```

	last_evaluation	number_project	average_monthly_hours	\
3058	0.359375	0.2	0.485981	
5091	0.218750	0.4	0.602804	
12996	0.781250	0.2	0.584112	
6296	0.500000	0.6	0.308411	
7584	0.968750	0.4	0.808411	

	time_spend_company	predicted_department	predicted_department_en	\
3058	0.125	7	sales	
5091	0.125	7	sales	
12996	0.125	7	sales	
6296	0.250	7	sales	
7584	0.125	7	sales	

	department	department_en
3058	9	technical
5091	8	support
12996	7	sales
6296	9	technical
7584	2	accounting

2.2 Hooray, we improved by approximately 0.01

2.3 Let's Visualize using a Residual Plot

```
In [142]: k=(reg.predict(X2_test) == y2_test) # Determine how many were predicted correctly
          k.value_counts()
```


```
Out[142]: False    4088
          True      412
          Name: sales, dtype: int64
```

```
In [143]: from sklearn.metrics import confusion_matrix
          cm=confusion_matrix(y2_test, reg.predict(X2_test), labels=y2_test.unique())
```

```
In [144]: plt.figure(figsize=(9,16))
          plot_confusion_matrix(cm,data['sales'].unique())
```

Confusion matrix, without normalization

```
[[ 0  0 143  0  0  0 142  0  0 548]
 [ 0  0 102  0  0  0 129  0  0 448]
 [ 0  0 218  0  0  0 219  0  0 789]
 [ 0  0  41  0  0  0  46  0  0 145]
 [ 0  0  40  0  0  0  42  0  0 148]
 [ 0  0  62  0  0  0  57  0  0 242]
 [ 0  0  35  0  0  0  51  0  0 155]
 [ 0  0  42  0  0  0  54  0  0 169]
 [ 0  0  31  0  0  0  38  0  0 111]
 [ 0  0  52  0  0  0  58  0  0 143]]
```



2.4 Did you try changing the amount of hidden layers??