

Week 5 - Neural Networks - Estimation - Python

October 10, 2018

1 Data Warehousing and Data Mining

1.1 Labs

1.1.1 Prepared by Gilroy Gordon

Contact Information SCIT ext. 3643

ggordon@utech@gmail.com

gilroy.gordon@utech.edu.jm

1.1.2 Week 5 - Neural Networks in Python

Additional Reference Resources:

http://scikit-learn.org/stable/modules/neural_networks_supervised.html

1.2 Objectives

- > Data Preprocessing
 - > Min Max Scaling
- > Data Transformation
- > Data Mining
 - > Neural Networks (Classification and Estimation)
- > Model Evaluation and Prediction
 - > Train/Test Split - 70/30
- > Presentation
 - > Plots

1.3 Import required libraries and acquire data

```
In [1]: # import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: data_path = './data/hr_data.csv' # Path to data file
data = pd.read_csv(data_path)
data.head(15)
```

```
Out[2]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	\
0	0.38	0.53	2	157	
1	0.80	0.86	5	262	
2	0.11	0.88	7	272	
3	0.72	0.87	5	223	
4	0.37	0.52	2	159	
5	0.41	0.50	2	153	
6	0.10	0.77	6	247	
7	0.92	0.85	5	259	
8	0.89	1.00	5	224	
9	0.42	0.53	2	142	
10	0.45	0.54	2	135	
11	0.11	0.81	6	305	
12	0.84	0.92	4	234	
13	0.41	0.55	2	148	
14	0.36	0.56	2	137	

	time_spend_company	Work_accident	left	promotion_last_5years	sales	\
0	3	0	1	0	sales	
1	6	0	1	0	sales	
2	4	0	1	0	sales	
3	5	0	1	0	sales	
4	3	0	1	0	sales	
5	3	0	1	0	sales	
6	4	0	1	0	sales	
7	5	0	1	0	sales	
8	5	0	1	0	sales	
9	3	0	1	0	sales	
10	3	0	1	0	sales	
11	4	0	1	0	sales	
12	5	0	1	0	sales	
13	3	0	1	0	sales	
14	3	0	1	0	sales	

	salary
0	low
1	medium
2	medium
3	low
4	low
5	low
6	low
7	low
8	low

```

9      low
10     low
11     low
12     low
13     low
14     low

```

```
In [3]: # What columns are in the data set ? Do they have spaces that I should consider
data.columns
```

```
Out[3]: Index(['satisfaction_level', 'last_evaluation', 'number_project',
              'average_monthly_hours', 'time_spend_company', 'Work_accident', 'left',
              'promotion_last_5years', 'sales', 'salary'],
              dtype='object')
```

1.4 Aim: Can we determine a person's Satisfaction Level based on the other factors?

satisfaction_level = a(last_evaluation) + b(number_project) + c(average_monthly_hours) + d(time_spend_company)

The coefficients a-d, what are they? What is the relationship between the variables? Does multicollinearity exist?

I have created a function below `create_label_encoder_dict` to assist with this. The function accepts a dataframe object and uses the `LabelEncoder` class from `sklearn.preprocessing` to encode (dummy encoding) or transform non-numerical columns to numbers. Finally it returns a dictionary object of all the encoders created for each column.

The `LabelEncoder` is a useful resource as it not only automatically transforms all values in a column but also keeps a track of what values were transformed from. i.e. It will change all Female to 0 and all Male to 1

```
In [4]: def create_label_encoder_dict(df):
        from sklearn.preprocessing import LabelEncoder

        label_encoder_dict = {}
        for column in df.columns:
            # Only create encoder for categorical data types
            if not np.issubdtype(df[column].dtype, np.number) and column != 'Age':
                label_encoder_dict[column] = LabelEncoder().fit(df[column])
        return label_encoder_dict

In [5]: label_encoders = create_label_encoder_dict(data)
        print("Encoded Values for each Label")
        print("="*32)
        for column in label_encoders:
            print("="*32)
            print('Encoder(%s) = %s' % (column, label_encoders[column].classes_ ))
            print(pd.DataFrame([range(0, len(label_encoders[column].classes_))], columns=label_en
```

Encoded Values for each Label

=====

=====

```
Encoder(sales) = ['IT' 'RandD' 'accounting' 'hr' 'management' 'marketing' 'product_mng'
                  'sales' 'support' 'technical']
```

Encoded Values

IT	0
RandD	1
accounting	2
hr	3
management	4
marketing	5
product_mng	6
sales	7
support	8
technical	9

=====

```
Encoder(salary) = ['high' 'low' 'medium']
```

Encoded Values

high	0
low	1
medium	2

```
In [6]: # Apply each encoder to the data set to obtain transformed values
data2 = data.copy() # create copy of initial data set
for column in data2.columns:
    if column in label_encoders:
        data2[column] = label_encoders[column].transform(data2[column])

print("Transformed data set")
print("="*32)
data2.head(15)
```

Transformed data set

=====

```
Out[6]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	\
0	0.38	0.53	2	157	
1	0.80	0.86	5	262	
2	0.11	0.88	7	272	
3	0.72	0.87	5	223	
4	0.37	0.52	2	159	
5	0.41	0.50	2	153	
6	0.10	0.77	6	247	
7	0.92	0.85	5	259	
8	0.89	1.00	5	224	

9	0.42	0.53	2	142
10	0.45	0.54	2	135
11	0.11	0.81	6	305
12	0.84	0.92	4	234
13	0.41	0.55	2	148
14	0.36	0.56	2	137

	time_spend_company	Work_accident	left	promotion_last_5years	sales	\
0	3	0	1	0	7	
1	6	0	1	0	7	
2	4	0	1	0	7	
3	5	0	1	0	7	
4	3	0	1	0	7	
5	3	0	1	0	7	
6	4	0	1	0	7	
7	5	0	1	0	7	
8	5	0	1	0	7	
9	3	0	1	0	7	
10	3	0	1	0	7	
11	4	0	1	0	7	
12	5	0	1	0	7	
13	3	0	1	0	7	
14	3	0	1	0	7	

	salary
0	1
1	2
2	2
3	1
4	1
5	1
6	1
7	1
8	1
9	1
10	1
11	1
12	1
13	1
14	1

```
In [7]: # separate our data into dependent (Y) and independent(X) variables
X_data = data2[['last_evaluation', 'number_project', 'average_monthly_hours', 'time_spend_co
Y_data = data2['satisfaction_level']
```

1.5 70/30 Train Test Split

We will split the data using a 70/30 split. i.e. 70% of the data will be randomly chosen to train the model and 30% will be used to evaluate the model

```
In [8]: from sklearn.model_selection import train_test_split
        X_train, X_test, y_train, y_test = train_test_split(X_data, Y_data, test_size=0.30)
```

```
In [9]: from sklearn.neural_network import MLPClassifier, MLPRegressor
```

```
In [10]: # Create an instance of linear regression
         reg = MLPRegressor()
```

```
In [11]: reg.fit(X_train,y_train)
```

```
Out[11]: MLPRegressor(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,
                      beta_2=0.999, early_stopping=False, epsilon=1e-08,
                      hidden_layer_sizes=(100,), learning_rate='constant',
                      learning_rate_init=0.001, max_iter=200, momentum=0.9,
                      nesterovs_momentum=True, power_t=0.5, random_state=None,
                      shuffle=True, solver='adam', tol=0.0001, validation_fraction=0.1,
                      verbose=False, warm_start=False)
```

```
In [12]: help(MLPRegressor)
```

Help on class MLPRegressor in module sklearn.neural_network.multilayer_perceptron:

```
class MLPRegressor(BaseMultilayerPerceptron, sklearn.base.RegressorMixin)
|   Multi-layer Perceptron regressor.
|
|   This model optimizes the squared-loss using LBFGS or stochastic gradient
|   descent.
|
|   .. versionadded:: 0.18
|
|   Parameters
|   -----
|   hidden_layer_sizes : tuple, length = n_layers - 2, default (100,)
|       The ith element represents the number of neurons in the ith
|       hidden layer.
|
|   activation : {'identity', 'logistic', 'tanh', 'relu'}, default 'relu'
|       Activation function for the hidden layer.
|
|       - 'identity', no-op activation, useful to implement linear bottleneck,
|         returns  $f(x) = x$ 
|
|       - 'logistic', the logistic sigmoid function,
|         returns  $f(x) = 1 / (1 + \exp(-x))$ .
```

- 'tanh', the hyperbolic tan function,
returns $f(x) = \tanh(x)$.
- 'relu', the rectified linear unit function,
returns $f(x) = \max(0, x)$

solver : {'lbfgs', 'sgd', 'adam'}, default 'adam'
The solver for weight optimization.

- 'lbfgs' is an optimizer in the family of quasi-Newton methods.
- 'sgd' refers to stochastic gradient descent.
- 'adam' refers to a stochastic gradient-based optimizer proposed by Kingma, Diederik, and Jimmy Ba

Note: The default solver 'adam' works pretty well on relatively large datasets (with thousands of training samples or more) in terms of both training time and validation score.
For small datasets, however, 'lbfgs' can converge faster and perform better.

alpha : float, optional, default 0.0001
L2 penalty (regularization term) parameter.

batch_size : int, optional, default 'auto'
Size of minibatches for stochastic optimizers.
If the solver is 'lbfgs', the classifier will not use minibatch.
When set to "auto", `batch_size=min(200, n_samples)`

learning_rate : {'constant', 'invscaling', 'adaptive'}, default 'constant'
Learning rate schedule for weight updates.

- 'constant' is a constant learning rate given by 'learning_rate_init'.
- 'invscaling' gradually decreases the learning rate ``learning_rate`` at each time step 't' using an inverse scaling exponent of 'power_t'.
$$\text{effective_learning_rate} = \text{learning_rate_init} / \text{pow}(t, \text{power_t})$$
- 'adaptive' keeps the learning rate constant to 'learning_rate_init' as long as training loss keeps decreasing.
Each time two consecutive epochs fail to decrease training loss by at least tol, or fail to increase validation score by at least tol if 'early_stopping' is on, the current learning rate is divided by 5.

Only used when solver='sgd'.

`learning_rate_init` : double, optional, default 0.001
 The initial learning rate used. It controls the step-size in updating the weights. Only used when `solver='sgd'` or `'adam'`.

`power_t` : double, optional, default 0.5
 The exponent for inverse scaling learning rate. It is used in updating effective learning rate when the `learning_rate` is set to `'invscaling'`. Only used when `solver='sgd'`.

`max_iter` : int, optional, default 200
 Maximum number of iterations. The solver iterates until convergence (determined by `'tol'`) or this number of iterations. For stochastic solvers (`'sgd'`, `'adam'`), note that this determines the number of epochs (how many times each data point will be used), not the number of gradient steps.

`shuffle` : bool, optional, default True
 Whether to shuffle samples in each iteration. Only used when `solver='sgd'` or `'adam'`.

`random_state` : int, RandomState instance or None, optional, default None
 If int, `random_state` is the seed used by the random number generator; If RandomState instance, `random_state` is the random number generator; If None, the random number generator is the RandomState instance used by ``np.random``.

`tol` : float, optional, default 1e-4
 Tolerance for the optimization. When the loss or score is not improving by at least `tol` for two consecutive iterations, unless ``learning_rate`` is set to `'adaptive'`, convergence is considered to be reached and training stops.

`verbose` : bool, optional, default False
 Whether to print progress messages to stdout.

`warm_start` : bool, optional, default False
 When set to True, reuse the solution of the previous call to fit as initialization, otherwise, just erase the previous solution.

`momentum` : float, default 0.9
 Momentum for gradient descent update. Should be between 0 and 1. Only used when `solver='sgd'`.

`nesterovs_momentum` : boolean, default True
 Whether to use Nesterov's momentum. Only used when `solver='sgd'` and `momentum > 0`.


```

| early_stopping : bool, default False
|     Whether to use early stopping to terminate training when validation
|     score is not improving. If set to true, it will automatically set
|     aside 10% of training data as validation and terminate training when
|     validation score is not improving by at least tol for two consecutive
|     epochs.
|     Only effective when solver='sgd' or 'adam'
|
| validation_fraction : float, optional, default 0.1
|     The proportion of training data to set aside as validation set for
|     early stopping. Must be between 0 and 1.
|     Only used if early_stopping is True
|
| beta_1 : float, optional, default 0.9
|     Exponential decay rate for estimates of first moment vector in adam,
|     should be in [0, 1). Only used when solver='adam'
|
| beta_2 : float, optional, default 0.999
|     Exponential decay rate for estimates of second moment vector in adam,
|     should be in [0, 1). Only used when solver='adam'
|
| epsilon : float, optional, default 1e-8
|     Value for numerical stability in adam. Only used when solver='adam'
|
| Attributes
| -----
| loss_ : float
|     The current loss computed with the loss function.
|
| coefs_ : list, length n_layers - 1
|     The ith element in the list represents the weight matrix corresponding
|     to layer i.
|
| intercepts_ : list, length n_layers - 1
|     The ith element in the list represents the bias vector corresponding to
|     layer i + 1.
|
| n_iter_ : int,
|     The number of iterations the solver has ran.
|
| n_layers_ : int
|     Number of layers.
|
| n_outputs_ : int
|     Number of outputs.
|
| out_activation_ : string

```

```

|         Name of the output activation function.
|
| Notes
| -----
| MLPRegressor trains iteratively since at each time step
| the partial derivatives of the loss function with respect to the model
| parameters are computed to update the parameters.
|
| It can also have a regularization term added to the loss function
| that shrinks model parameters to prevent overfitting.
|
| This implementation works with data represented as dense and sparse numpy
| arrays of floating point values.
|
| References
| -----
| Hinton, Geoffrey E.
|     "Connectionist learning procedures." Artificial intelligence 40.1
|     (1989): 185-234.
|
| Glorot, Xavier, and Yoshua Bengio. "Understanding the difficulty of
|     training deep feedforward neural networks." International Conference
|     on Artificial Intelligence and Statistics. 2010.
|
| He, Kaiming, et al. "Delving deep into rectifiers: Surpassing human-level
|     performance on imagenet classification." arXiv preprint
|     arXiv:1502.01852 (2015).
|
| Kingma, Diederik, and Jimmy Ba. "Adam: A method for stochastic
|     optimization." arXiv preprint arXiv:1412.6980 (2014).
|
| Method resolution order:
|     MLPRegressor
|     BaseMultilayerPerceptron
|     abc.NewBase
|     sklearn.base.BaseEstimator
|     sklearn.base.RegressorMixin
|     builtins.object
|
| Methods defined here:
|
|     __init__(self, hidden_layer_sizes=(100,), activation='relu', solver='adam', alpha=0.0001, ba
|         Initialize self. See help(type(self)) for accurate signature.
|
|     predict(self, X)
|         Predict using the multi-layer perceptron model.
|
|     Parameters

```

```

|         -----
|         X : {array-like, sparse matrix}, shape (n_samples, n_features)
|             The input data.
|
|         Returns
|         -----
|         y : array-like, shape (n_samples, n_outputs)
|             The predicted values.
|
|         -----
|         Data and other attributes defined here:
|
|         __abstractmethods__ = frozenset()
|
|         -----
|         Methods inherited from BaseMultilayerPerceptron:
|
|         fit(self, X, y)
|             Fit the model to data matrix X and target(s) y.
|
|         Parameters
|         -----
|         X : array-like or sparse matrix, shape (n_samples, n_features)
|             The input data.
|
|         y : array-like, shape (n_samples,) or (n_samples, n_outputs)
|             The target values (class labels in classification, real numbers in
|             regression).
|
|         Returns
|         -----
|         self : returns a trained MLP model.
|
|         -----
|         Data descriptors inherited from BaseMultilayerPerceptron:
|
|         partial_fit
|             Fit the model to data matrix X and target y.
|
|         Parameters
|         -----
|         X : {array-like, sparse matrix}, shape (n_samples, n_features)
|             The input data.
|
|         y : array-like, shape (n_samples,)
|             The target values.
|
|         Returns

```

```

|         -----
|         self : returns a trained MLP model.
|
| -----
| Methods inherited from sklearn.base.BaseEstimator:
|
|     __getstate__(self)
|
|     __repr__(self)
|         Return repr(self).
|
|     __setstate__(self, state)
|
|     get_params(self, deep=True)
|         Get parameters for this estimator.
|
|         Parameters
|         -----
|         deep : boolean, optional
|             If True, will return the parameters for this estimator and
|             contained subobjects that are estimators.
|
|         Returns
|         -----
|         params : mapping of string to any
|             Parameter names mapped to their values.
|
|     set_params(self, **params)
|         Set the parameters of this estimator.
|
|         The method works on simple estimators as well as on nested objects
|         (such as pipelines). The latter have parameters of the form
|         ``<component>__<parameter>`` so that it's possible to update each
|         component of a nested object.
|
|         Returns
|         -----
|         self
|
| -----
| Data descriptors inherited from sklearn.base.BaseEstimator:
|
|     __dict__
|         dictionary for instance variables (if defined)
|
|     __weakref__
|         list of weak references to the object (if defined)

```

```

| -----
| Methods inherited from sklearn.base.RegressorMixin:
|
| score(self, X, y, sample_weight=None)
|     Returns the coefficient of determination R^2 of the prediction.
|
|     The coefficient R^2 is defined as (1 - u/v), where u is the residual
|     sum of squares ((y_true - y_pred) ** 2).sum() and v is the total
|     sum of squares ((y_true - y_true.mean()) ** 2).sum().
|     The best possible score is 1.0 and it can be negative (because the
|     model can be arbitrarily worse). A constant model that always
|     predicts the expected value of y, disregarding the input features,
|     would get a R^2 score of 0.0.
|
|     Parameters
|     -----
|     X : array-like, shape = (n_samples, n_features)
|         Test samples.
|
|     y : array-like, shape = (n_samples) or (n_samples, n_outputs)
|         True values for X.
|
|     sample_weight : array-like, shape = [n_samples], optional
|         Sample weights.
|
|     Returns
|     -----
|     score : float
|         R^2 of self.predict(X) wrt. y.

```

```
In [18]: reg.n_layers_ # Number of layers utilized
```

```
Out[18]: 3
```

```
In [20]: # Make predictions using the testing set
         test_predicted = reg.predict(X_test)
         test_predicted
```

```
Out[20]: array([0.66908448, 0.65429664, 0.60327367, ..., 0.47367503, 0.72666922,
                0.57838912])
```

```
In [21]: data3 = X_test.copy()
         data3['predicted_satisfaction_level']=test_predicted
         data3['satisfaction_level']=y_test
         data3.head()
```

```
Out[21]:      last_evaluation  number_project  average_monthly_hours  \
6335                0.81                5                143
```

7020	0.73	3	138
8864	0.56	3	142
7454	0.57	4	141
13361	0.61	4	118

	time_spend_company	predicted_satisfaction_level	satisfaction_level
6335	2	0.669084	0.80
7020	3	0.654297	0.62
8864	3	0.603274	0.81
7454	3	0.577754	0.49
13361	5	0.541316	0.54

```
In [23]: from sklearn.metrics import mean_squared_error, r2_score
```

```
In [24]: # The mean squared error don't worry guys we can do this
print("Mean squared error: %.2f" % mean_squared_error(y_test, test_predicted))
```

Mean squared error: 0.06

```
In [25]: # Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % r2_score(y_test, test_predicted))
```

Variance score: 0.07

2 Preprocessing

```
In [30]: def create_min_max_scaler_dict(df):
    from sklearn.preprocessing import MinMaxScaler
    min_max_scaler_dict = {}
    for column in df.columns:
        # Only create encoder for categorical data types
        if np.issubdtype(df[column].dtype, np.number):
            min_max_scaler_dict[column] = MinMaxScaler().fit(pd.DataFrame(df[column]))
    return min_max_scaler_dict
```

```
In [31]: min_max_scalers = create_min_max_scaler_dict(data)
print("Min Max Values for each Label")
print("="*32)
min_max_scalers
```

Min Max Values for each Label
=====

```
Out[31]: {'Work_accident': MinMaxScaler(copy=True, feature_range=(0, 1)),
'average_montly_hours': MinMaxScaler(copy=True, feature_range=(0, 1)),
'last_evaluation': MinMaxScaler(copy=True, feature_range=(0, 1)),
```

```

    'left': MinMaxScaler(copy=True, feature_range=(0, 1)),
    'number_project': MinMaxScaler(copy=True, feature_range=(0, 1)),
    'promotion_last_5years': MinMaxScaler(copy=True, feature_range=(0, 1)),
    'satisfaction_level': MinMaxScaler(copy=True, feature_range=(0, 1)),
    'time_spend_company': MinMaxScaler(copy=True, feature_range=(0, 1))}

In [48]: #retrieving a scacler
         time_spend_company_scaler=min_max_scalers['time_spend_company']

In [49]: time_spend_company_scaler

Out[49]: MinMaxScaler(copy=True, feature_range=(0, 1))

In [50]: time_spend_company_scaler.data_max_ #Maximum value

Out[50]: array([10.])

In [51]: time_spend_company_scaler.data_min_ # Minimum value

Out[51]: array([2.])

In [52]: time_spend_company_scaler.data_range_ # Range = Max- Min

Out[52]: array([8.])

In [53]: pd.DataFrame([
        {
            'column':col,
            'min':min_max_scalers[col].data_min_[0],
            'max':min_max_scalers[col].data_max_[0],
            'range':min_max_scalers[col].data_range_[0] } for col in min_max_scalers])

Out[53]:
         column  max  min  range
0  promotion_last_5years    1  0.00   1.00
1    time_spend_company   10  2.00   8.00
2      number_project    7  2.00   5.00
3              left    1  0.00   1.00
4    last_evaluation    1  0.36   0.64
5  average_monthly_hours  310 96.00  214.00
6    satisfaction_level    1  0.09   0.91
7        Work_accident    1  0.00   1.00

In [55]: # Apply each scaler to the data set to obtain transformed values
         data3 = data2.copy() # create copy of initial data set
         for column in data3.columns:
             if column in min_max_scalers:
                 data3[column] = min_max_scalers[column].transform(pd.DataFrame(data3[column]))

         print("Transformed data set")
         print("="*32)
         data3.head(15)

```

Transformed data set

=====

```
Out[55]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	\
0	0.318681	0.265625	0.0	0.285047	
1	0.780220	0.781250	0.6	0.775701	
2	0.021978	0.812500	1.0	0.822430	
3	0.692308	0.796875	0.6	0.593458	
4	0.307692	0.250000	0.0	0.294393	
5	0.351648	0.218750	0.0	0.266355	
6	0.010989	0.640625	0.8	0.705607	
7	0.912088	0.765625	0.6	0.761682	
8	0.879121	1.000000	0.6	0.598131	
9	0.362637	0.265625	0.0	0.214953	
10	0.395604	0.281250	0.0	0.182243	
11	0.021978	0.703125	0.8	0.976636	
12	0.824176	0.875000	0.4	0.644860	
13	0.351648	0.296875	0.0	0.242991	
14	0.296703	0.312500	0.0	0.191589	

	time_spend_company	Work_accident	left	promotion_last_5years	sales	\
0	0.125	0	1	0	7	
1	0.500	0	1	0	7	
2	0.250	0	1	0	7	
3	0.375	0	1	0	7	
4	0.125	0	1	0	7	
5	0.125	0	1	0	7	
6	0.250	0	1	0	7	
7	0.375	0	1	0	7	
8	0.375	0	1	0	7	
9	0.125	0	1	0	7	
10	0.125	0	1	0	7	
11	0.250	0	1	0	7	
12	0.375	0	1	0	7	
13	0.125	0	1	0	7	
14	0.125	0	1	0	7	

	salary
0	1
1	2
2	2
3	1
4	1
5	1
6	1
7	1
8	1


```

9         1
10        1
11        1
12        1
13        1
14        1

```

```
In [56]: data3.describe()
```

```

Out[56]:      satisfaction_level  last_evaluation  number_project  \
count      14999.000000      14999.000000      14999.000000
mean         0.574542         0.556409         0.360611
std          0.273220         0.267452         0.246518
min          0.000000         0.000000         0.000000
25%          0.384615         0.312500         0.200000
50%          0.604396         0.562500         0.400000
75%          0.802198         0.796875         0.600000
max          1.000000         1.000000         1.000000

      average_monthly_hours  time_spend_company  Work_accident  left  \
count      14999.000000      14999.000000      14999.000000  14999.000000
mean         0.490889         0.187279         0.144610      0.238083
std          0.233379         0.182517         0.351719      0.425924
min          0.000000         0.000000         0.000000      0.000000
25%          0.280374         0.125000         0.000000      0.000000
50%          0.485981         0.125000         0.000000      0.000000
75%          0.696262         0.250000         0.000000      0.000000
max          1.000000         1.000000         1.000000      1.000000

      promotion_last_5years      sales      salary
count      14999.000000  14999.000000  14999.000000
mean         0.021268      5.870525      1.347290
std          0.144281      2.868786      0.625819
min          0.000000      0.000000      0.000000
25%          0.000000      4.000000      1.000000
50%          0.000000      7.000000      1.000000
75%          0.000000      8.000000      2.000000
max          1.000000      9.000000      2.000000

```

```
In [57]: # separate our data into dependent (Y) and independent(X) variables
```

```

X2_data = data3[['last_evaluation', 'number_project', 'average_monthly_hours', 'time_spend_company']]
Y2_data = data3['satisfaction_level']

```

2.1 70/30 Train Test Split

We will split the data using a 70/30 split. i.e. 70% of the data will be randomly chosen to train the model and 30% will be used to evaluate the model

```

In [58]: from sklearn.model_selection import train_test_split
X2_train, X2_test, y2_train, y2_test = train_test_split(X2_data, Y2_data, test_size=0.3)

```

```

In [60]: # Create an instance of linear regression
reg2 = MLPRegressor()
reg2.fit(X2_train,y2_train)

Out[60]: MLPRegressor(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,
    beta_2=0.999, early_stopping=False, epsilon=1e-08,
    hidden_layer_sizes=(100,), learning_rate='constant',
    learning_rate_init=0.001, max_iter=200, momentum=0.9,
    nesterovs_momentum=True, power_t=0.5, random_state=None,
    shuffle=True, solver='adam', tol=0.0001, validation_fraction=0.1,
    verbose=False, warm_start=False)

In [61]: reg2.n_layers_

Out[61]: 3

In [63]: # Make predictions using the testing set
test2_predicted = reg2.predict(X2_test)
test2_predicted

Out[63]: array([0.57568627, 0.57893697, 0.5534623 , ..., 0.55752807, 0.71729991,
    0.58278402])

In [64]: # The mean squared error don't worry guys we can do this
print("Mean squared error: %.2f" % mean_squared_error(y2_test, test2_predicted))

Mean squared error: 0.05

```

2.2 Hooray, we improved by approximately 0.01

2.3 Let's Visualize using a Residual Plot

```

In [65]: import seaborn as sns
sns.set(style="whitegrid")

In [71]: sns.residplot(y2_test, test2_predicted, color="g")
plt.title("Residual Plot")
plt.ylabel("Error")

Out[71]: <matplotlib.text.Text at 0x7f1a7da71518>

```



2.4 Did you try changing the amount of hidden layers??