

Week 6 - Clustering - Python

October 24, 2018

1 Data Warehousing and Data Mining

1.1 Labs

1.1.1 Prepared by Gilroy Gordon

Contact Information SCIT ext. 3643

ggordonutech@gmail.com

gilroy.gordon@utech.edu.jm

1.1.2 Week 6 - Clustering in Python

Additional Reference Resources:

<http://scikit-learn.org/stable/modules/clustering.html>

1.2 Objectives

- > Data Preprocessing
 - > Missing Values (Na and nulls)
- > Data Mining
 - > Clustering (Kmeans)
- > Visualizations
 - > Elbow method (Within Cluster Sum of Squares)
- > Discussion on how to proceed

1.3 Aim: Am I able to segment groups based on

1.4 Import required libraries and acquire data

```
In [1]: # import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: data_path = './data/credit-card-data.csv' # Path to data file
data = pd.read_csv(data_path)
data.head(15)
```

```
Out[2]:
```

	cust_id	balance	balance_frequency	purchases	oneoff_purchases \
0	C10001	40.900749	0.818182	95.40	0.00
1	C10002	3202.467416	0.909091	0.00	0.00
2	C10003	2495.148862	1.000000	773.17	773.17
3	C10004	1666.670542	0.636364	1499.00	1499.00
4	C10005	817.714335	1.000000	16.00	16.00
5	C10006	1809.828751	1.000000	1333.28	0.00
6	C10007	627.260806	1.000000	7091.01	6402.63
7	C10008	1823.652743	1.000000	436.20	0.00
8	C10009	1014.926473	1.000000	861.49	661.49
9	C10010	152.225975	0.545455	1281.60	1281.60
10	C10011	1293.124939	1.000000	920.12	0.00
11	C10012	630.794744	0.818182	1492.18	1492.18
12	C10013	1516.928620	1.000000	3217.99	2500.23
13	C10014	921.693369	1.000000	2137.93	419.96
14	C10015	2772.772734	1.000000	0.00	0.00

	installments_purchases	cash_advance	purchases_frequency \
0	95.40	0.000000	0.166667
1	0.00	6442.945483	0.000000
2	0.00	0.000000	1.000000
3	0.00	205.788017	0.083333
4	0.00	0.000000	0.083333
5	1333.28	0.000000	0.666667
6	688.38	0.000000	1.000000
7	436.20	0.000000	1.000000
8	200.00	0.000000	0.333333
9	0.00	0.000000	0.166667
10	920.12	0.000000	1.000000
11	0.00	0.000000	0.250000
12	717.76	0.000000	1.000000
13	1717.97	0.000000	0.750000
14	0.00	346.811390	0.000000

	oneoff_purchases_frequency	purchases_installments_frequency \
0	0.000000	0.083333
1	0.000000	0.000000
2	1.000000	0.000000
3	0.083333	0.000000
4	0.083333	0.000000
5	0.000000	0.583333
6	1.000000	1.000000
7	0.000000	1.000000
8	0.083333	0.250000

9	0.166667	0.000000
10	0.000000	1.000000
11	0.250000	0.000000
12	0.250000	0.916667
13	0.166667	0.750000
14	0.000000	0.000000

	cash_advance_frequency	cash_advance_trx	purchases_trx	credit_limit \
0	0.000000	0	2	1000
1	0.250000	4	0	7000
2	0.000000	0	12	7500
3	0.083333	1	1	7500
4	0.000000	0	1	1200
5	0.000000	0	8	1800
6	0.000000	0	64	13500
7	0.000000	0	12	2300
8	0.000000	0	5	7000
9	0.000000	0	3	11000
10	0.000000	0	12	1200
11	0.000000	0	6	2000
12	0.000000	0	26	3000
13	0.000000	0	26	7500
14	0.083333	1	0	3000

	payments	minimum_payments	prc_full_payment	tenure
0	201.802084	139.509787	0.000000	12
1	4103.032597	1072.340217	0.222222	12
2	622.066742	627.284787	0.000000	12
3	0.000000	NaN	0.000000	12
4	678.334763	244.791237	0.000000	12
5	1400.057770	2407.246035	0.000000	12
6	6354.314328	198.065894	1.000000	12
7	679.065082	532.033990	0.000000	12
8	688.278568	311.963409	0.000000	12
9	1164.770591	100.302262	0.000000	12
10	1083.301007	2172.697765	0.000000	12
11	705.618627	155.549069	0.000000	12
12	608.263689	490.207013	0.250000	12
13	1655.891435	251.137986	0.083333	12
14	805.647974	989.962866	0.000000	12

In [3]: *# What columns are in the data set ? Do they have spaces that I should consider*
data.columns

Out[3]: Index(['cust_id', 'balance', 'balance_frequency', 'purchases',
'oneoff_purchases', 'installments_purchases', 'cash_advance',
'purchases_frequency', 'oneoff_purchases_frequency',
'purchases_installments_frequency', 'cash_advance_frequency',

```

        'cash_advance_trx', 'purchases_trx', 'credit_limit', 'payments',
        'minimum_payments', 'prc_full_payment', 'tenure'],
        dtype='object')

```

1.4.1 Exploring our data

Methods also available in Kaggle Bot default kernel. Some Modifications made for compatibility

```

In [4]: # Histogram of column data
def plotHistogram(df, nHistogramShown, nHistogramPerRow):
    nunique = len(df.columns)
    df = df[[col for col in df if nunique[col] > 1 and nunique[col] < 50]] # For display
    nRow, nCol = df.shape
    columnNames = list(df)
    nHistRow = (nCol + nHistogramPerRow - 1) / nHistogramPerRow
    plt.figure(num=None, figsize=(6*nHistogramPerRow, 8*nHistRow), dpi=80, facecolor='w')
    for i in range(min(nCol, nHistogramShown)):
        plt.subplot(nHistRow, nHistogramPerRow, i+1)
        df.iloc[:,i].hist()
        plt.ylabel('counts')
        plt.xticks(rotation=90)
        plt.title('{0} (column {1})'.format(columnNames[i],i))
    plt.tight_layout(pad=1.0, w_pad=1.0, h_pad=1.0)
    plt.show()

# Correlation matrix
def plotCorrelationMatrix(df, graphWidth):
    filename = df.dataframeName
    df = df.dropna('columns') # drop columns with NaN
    df = df[[col for col in df if df[col].nunique() > 1]] # keep columns where there are
    if df.shape[1] < 2:
        print('No correlation plots shown: The number of non-NaN or constant columns ({0}
        return
    corr = df.corr()
    plt.figure(num=None, figsize=(graphWidth, graphWidth), dpi=80, facecolor='w', edgecolor='k')
    corrMat = plt.matshow(corr, fignum = 1)
    plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
    plt.yticks(range(len(corr.columns)), corr.columns)
    plt.gca().xaxis.tick_bottom()
    plt.colorbar(corrMat)
    plt.title('Correlation Matrix for {0}'.format(filename), fontsize=15)
    plt.show()

# Scatter and density plots
def plotScatterMatrix(df, plotSize, textSize):
    df = df.select_dtypes(include=[np.number]) # keep only numerical columns
    # Remove rows and columns that would lead to df being singular

```

```

df = df.dropna('columns')
df = df[[col for col in df if df[col].nunique() > 1]] # keep columns where there are
columnNames = list(df)
if len(columnNames) > 10: # reduce the number of columns for matrix inversion of kern
    columnNames = columnNames[:10]
df = df[columnNames]
ax = pd.plotting.scatter_matrix(df, alpha=0.75, figsize=[plotSize, plotSize], diagonal=corr
corr = df.corr().values
for i, j in zip(*plt.np.triu_indices_from(ax, k = 1)):
    ax[i, j].annotate('Corr. coef = %.3f' % corr[i, j], (0.8, 0.2), xycoords='axes
plt.suptitle('Scatter and Density Plot')
plt.show()

```

In [5]: data.describe()

```

Out[5]:

```

	balance	balance_frequency	purchases	oneoff_purchases	\
count	8950.000000	8950.000000	8950.000000	8950.000000	
mean	1564.474828	0.877271	1003.204834	592.437371	
std	2081.531879	0.236904	2136.634782	1659.887917	
min	0.000000	0.000000	0.000000	0.000000	
25%	128.281915	0.888889	39.635000	0.000000	
50%	873.385231	1.000000	361.280000	38.000000	
75%	2054.140036	1.000000	1110.130000	577.405000	
max	19043.138560	1.000000	49039.570000	40761.250000	

	installments_purchases	cash_advance	purchases_frequency	\
count	8950.000000	8950.000000	8950.000000	
mean	411.067645	978.871112	0.490351	
std	904.338115	2097.163877	0.401371	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.083333	
50%	89.000000	0.000000	0.500000	
75%	468.637500	1113.821139	0.916667	
max	22500.000000	47137.211760	1.000000	


	oneoff_purchases_frequency	purchases_installments_frequency	\
count	8950.000000	8950.000000	
mean	0.202458	0.364437	
std	0.298336	0.397448	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	0.083333	0.166667	
75%	0.300000	0.750000	
max	1.000000	1.000000	


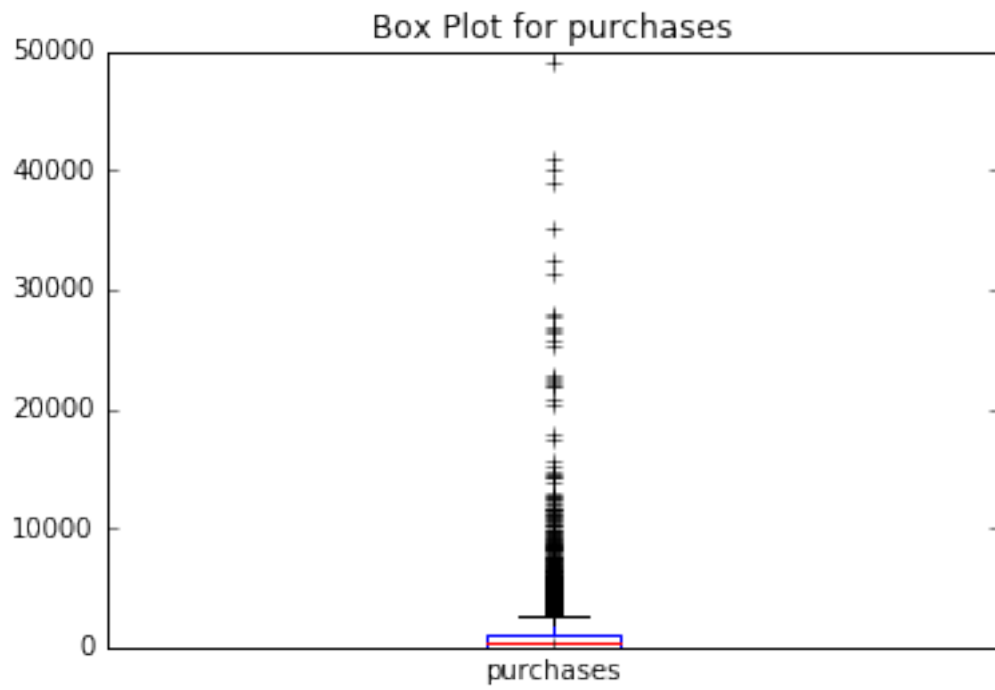
	cash_advance_frequency	cash_advance_trx	purchases_trx	credit_limit	\
count	8950.000000	8950.000000	8950.000000	8949.000000	
mean	0.135144	3.248827	14.709832	4494.449450	



std	0.200121	6.824647	24.857649	3638.815725
min	0.000000	0.000000	0.000000	50.000000
25%	0.000000	0.000000	1.000000	1600.000000
50%	0.000000	0.000000	7.000000	3000.000000
75%	0.222222	4.000000	17.000000	6500.000000
max	1.500000	123.000000	358.000000	30000.000000



	payments	minimum_payments	prc_full_payment	tenure
count	8950.000000	8637.000000	8950.000000	8950.000000
mean	1733.143852	864.206542	0.153715	11.517318
std	2895.063757	2372.446607	0.292499	1.338331
min	0.000000	0.019163	0.000000	6.000000
25%	383.276166	169.123707	0.000000	12.000000
50%	856.901546	312.343947	0.000000	12.000000
75%	1901.134317	825.485459	0.142857	12.000000
max	50721.483360	76406.207520	1.000000	12.000000



```
In [6]: # Let's view the distribution of the data, where is it possible to find groups?
# We are using boxplots of all the columns except the first (cust_id which is a string)
for col in data.columns[2:]:
    data[col].plot(kind='box')
    plt.title('Box Plot for '+col)
    plt.show()
```

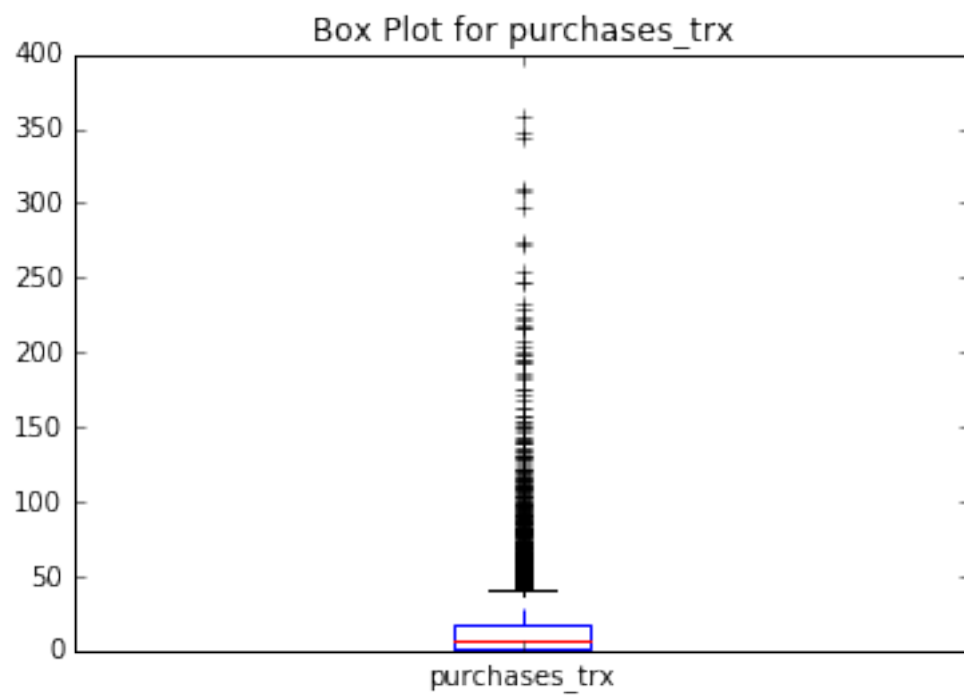



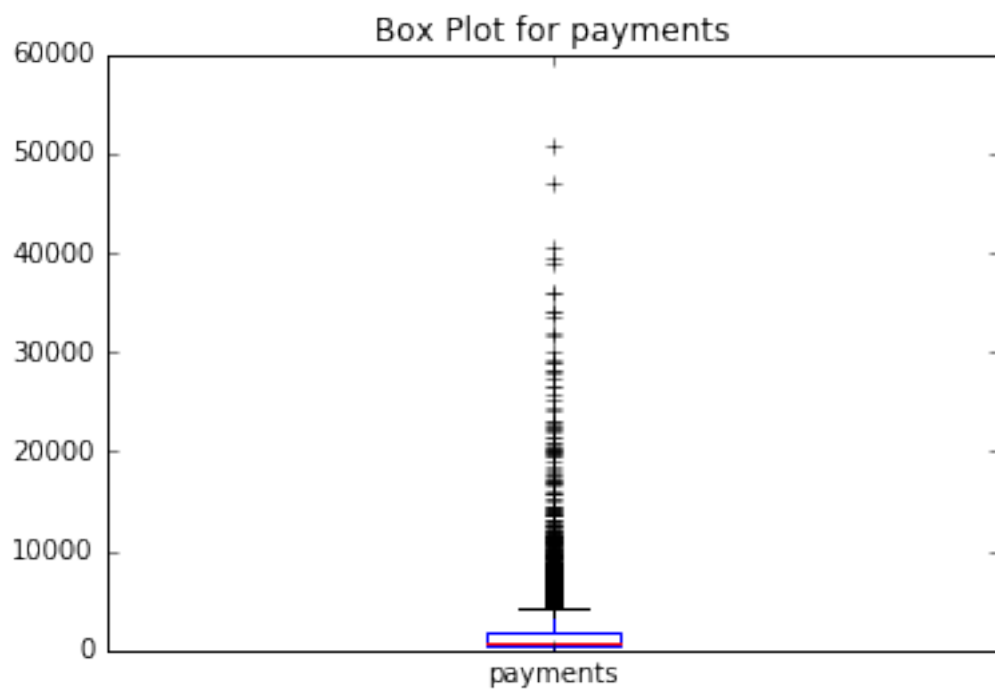




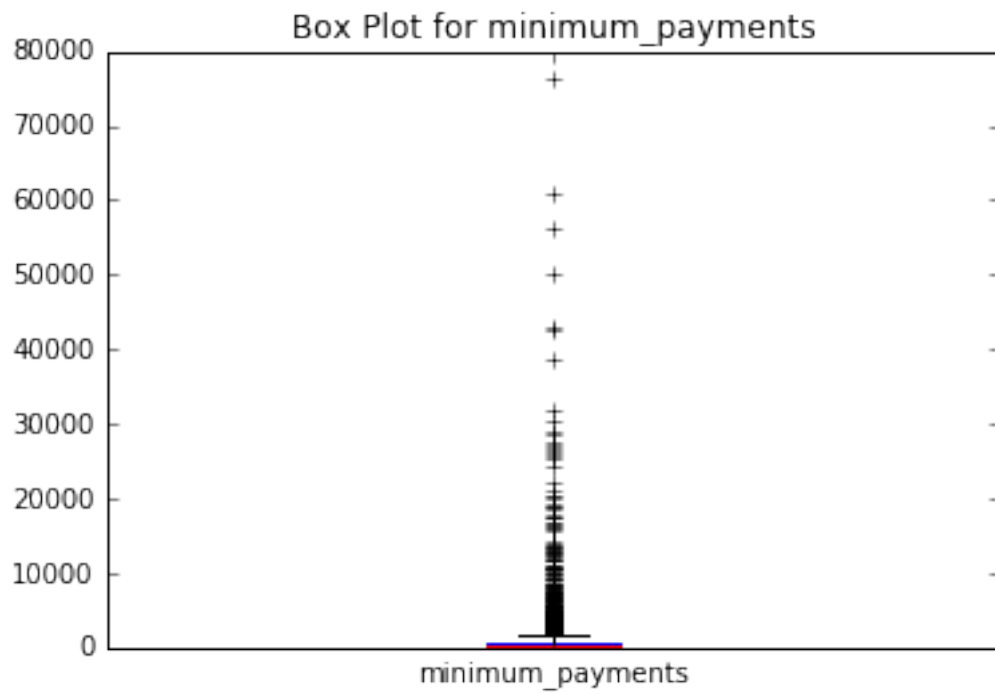


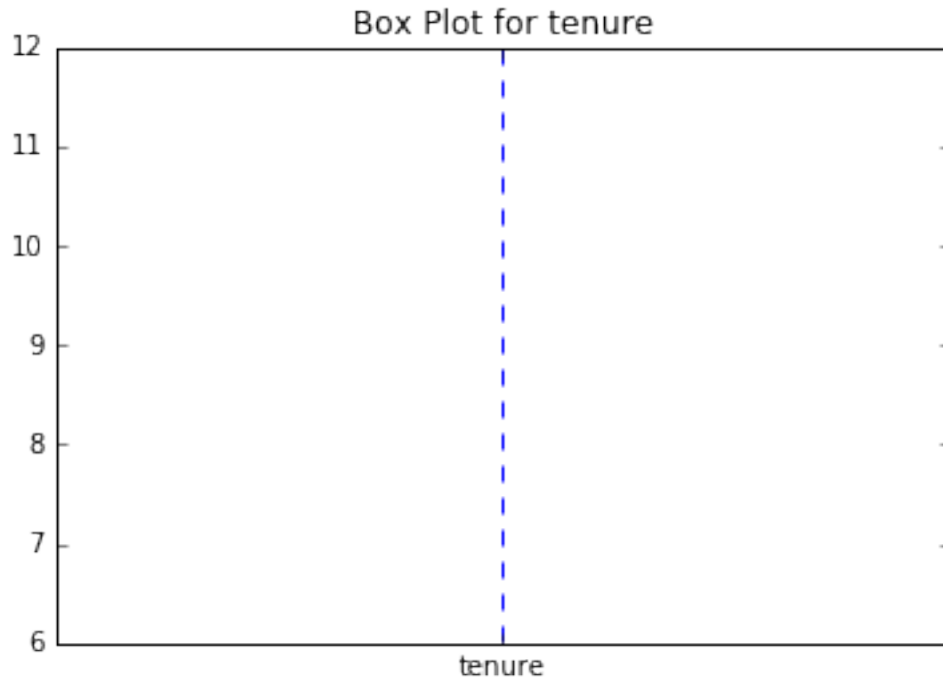












1.5 Aim: Can we identify groups based on purchases and payments?

If that is the case, we could offer different payment plans based on different purchases.


```
In [7]: cluster_data = data[['purchases', 'payments']]
        cluster_data.head()
```

```
Out[7]:
```

	purchases	payments
0	95.40	201.802084
1	0.00	4103.032597
2	773.17	622.066742
3	1499.00	0.000000
4	16.00	678.334763

```
In [8]: cluster_data.plot(kind='scatter', x='purchases', y='payments')
```

```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x7f836b17c4a8>
```



```
In [9]: # Is there any missing data
missing_data_results = cluster_data.isnull().sum()
print(missing_data_results)

# perform imputation with median values
# not require since none missing
#cluster_data = cluster_data.fillna( data.median() )
```

```
purchases    0
payments     0
dtype: int64
```

```
In [10]: #retrieve just the values for all columns except customer id
data_values = cluster_data.iloc[ :, 1:].values
data_values
```

```
Out[10]: array([[ 201.802084],
 [4103.032597],
 [ 622.066742],
 ...,
 [ 81.270775],
 [ 52.549959],
 [ 63.165404]])
```


```

In [11]: #import KMeans algorithm
         from sklearn.cluster import KMeans

In [12]: # Use the Elbow method to find a good number of clusters using WCSS (within-cluster sum
         wcss = []
         for i in range( 1, 15 ):
             kmeans = KMeans(n_clusters=i, init="k-means++", n_init=10, max_iter=300)
             kmeans.fit_predict( data_values )
             wcss.append( kmeans.inertia_ )

         plt.plot( wcss, 'ro-', label="WCSS")
         plt.title("Computing WCSS for KMeans++")
         plt.xlabel("Number of clusters")
         plt.ylabel("WCSS")
         plt.show()

```



We're seeing an elbow at approx 3, so let's try 3 groups

```

In [13]: kmeans = KMeans(n_clusters=3, init="k-means++", n_init=10, max_iter=300)
         cluster_data["cluster"] = kmeans.fit_predict( data_values )
         cluster_data

```

/usr/local/lib/python3.5/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>

```
Out[13]:
```

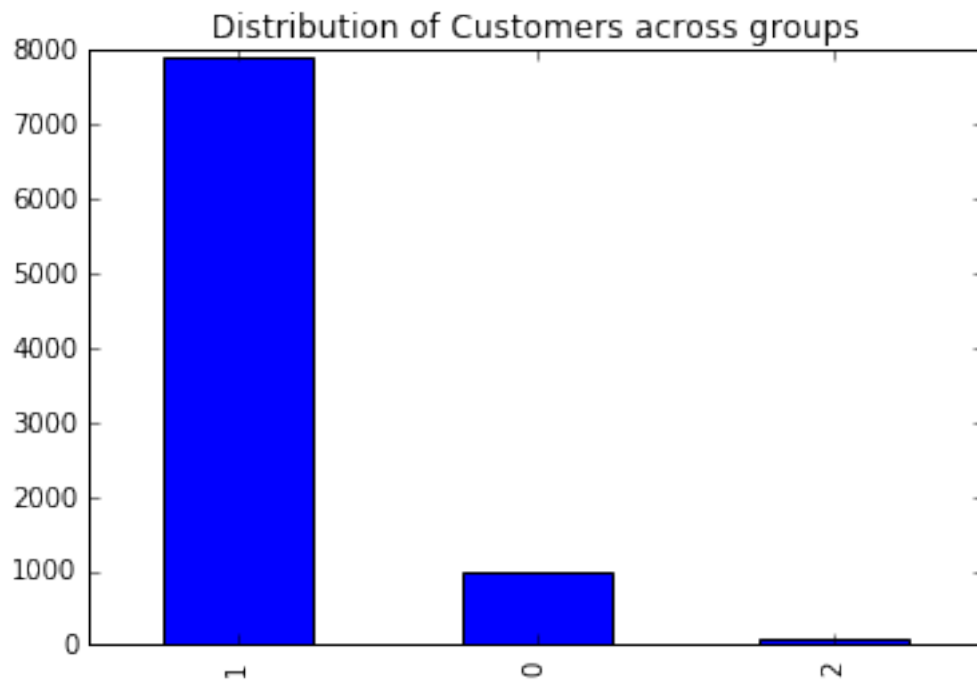
	purchases	payments	cluster
0	95.40	201.802084	1
1	0.00	4103.032597	0
2	773.17	622.066742	1
3	1499.00	0.000000	1
4	16.00	678.334763	1
5	1333.28	1400.057770	1
6	7091.01	6354.314328	0
7	436.20	679.065082	1
8	861.49	688.278568	1
9	1281.60	1164.770591	1
10	920.12	1083.301007	1
11	1492.18	705.618627	1
12	3217.99	608.263689	1
13	2137.93	1655.891435	1
14	0.00	805.647974	1
15	1611.70	1993.439277	1
16	0.00	391.974562	1
17	519.00	254.590662	1
18	504.35	1720.837373	1
19	398.64	1053.980464	1
20	176.68	223.068600	1
21	6359.95	2077.959051	1
22	815.90	2359.629958	1
23	4248.35	9479.043842	0
24	0.00	1422.726707	1
25	399.60	215.306142	1
26	102.00	890.178845	1
27	233.28	207.773715	1
28	387.05	1601.448347	1
29	100.00	160.767773	1
...
8920	0.00	54.795084	1
8921	57.42	68.462579	1
8922	145.98	53.676054	1
8923	1898.88	669.039640	1
8924	74.00	214.921009	1
8925	418.59	422.538988	1
8926	580.00	641.303466	1
8927	315.20	231.274641	1
8928	500.00	456.745027	1
8929	0.00	0.000000	1
8930	84.00	124.373736	1

8931	235.80	189.090274	1
8932	180.00	138.203240	1
8933	619.60	106.138603	1
8934	110.50	161.476789	1
8935	465.90	0.000000	1
8936	712.50	605.716356	1
8937	0.00	117.738787	1
8938	0.00	1397.770131	1
8939	734.40	72.530037	1
8940	591.24	475.523262	1
8941	214.55	966.202912	1
8942	113.28	94.488828	1
8943	20.90	58.644883	1
8944	1012.73	0.000000	1
8945	291.12	325.594462	1
8946	300.00	275.861322	1
8947	144.40	81.270775	1
8948	0.00	52.549959	1
8949	1093.25	63.165404	1

[8950 rows x 3 columns]

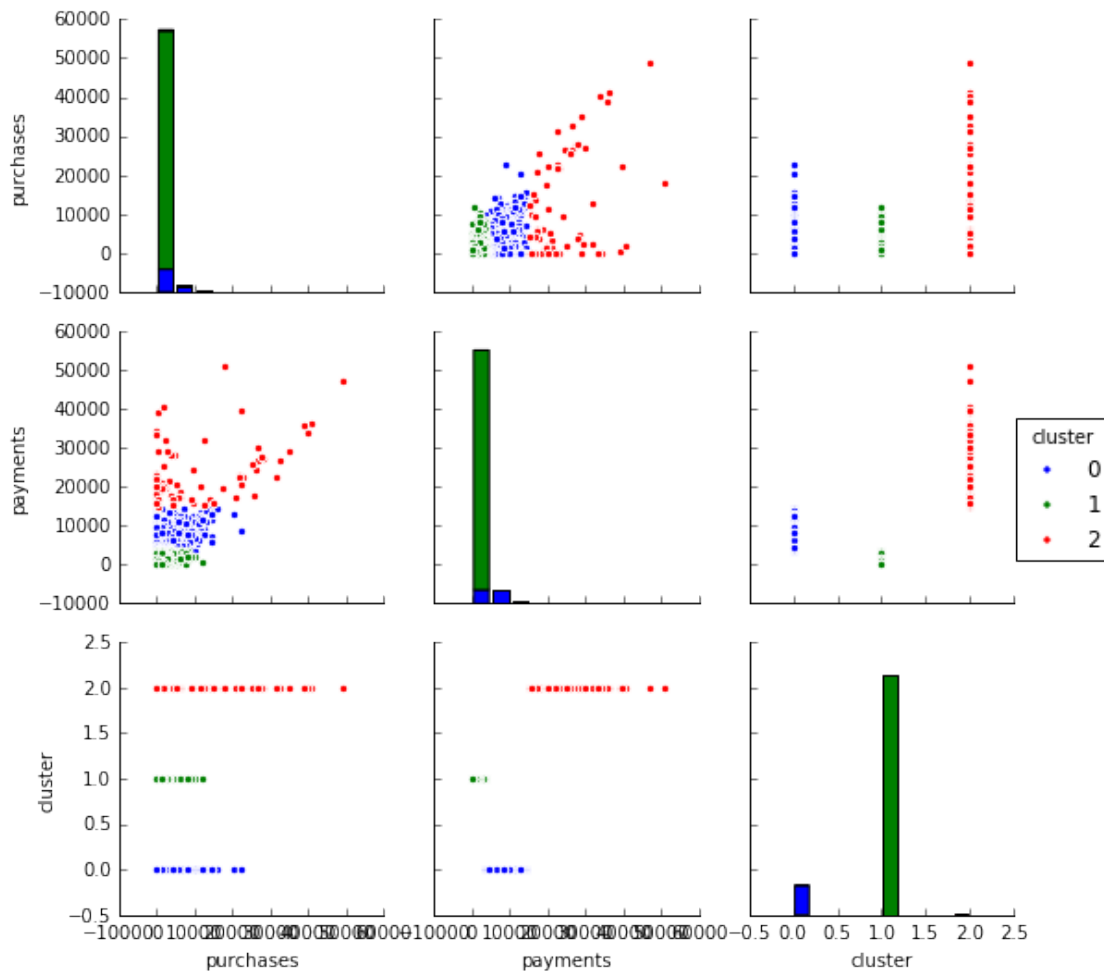
In [24]: `cluster_data['cluster'].value_counts().plot(kind='bar',title='Distribution of Customers`

`Out[24]: <matplotlib.axes._subplots.AxesSubplot at 0x7f83647e24a8>`



```
In [15]: sns.pairplot( cluster_data, hue="cluster")
```

```
Out[15]: <seaborn.axisgrid.PairGrid at 0x7f836dd89f98>
```



Looks nice, but what are we really seeing? Let's attempt to describe the groups

```
In [17]: grouped_cluster_data = cluster_data.groupby('cluster')
grouped_cluster_data
```

```
Out[17]: <pandas.core.groupby.DataFrameGroupBy object at 0x7f836531b128>
```

```
In [18]: grouped_cluster_data.describe()
```


```
Out[18]:
```


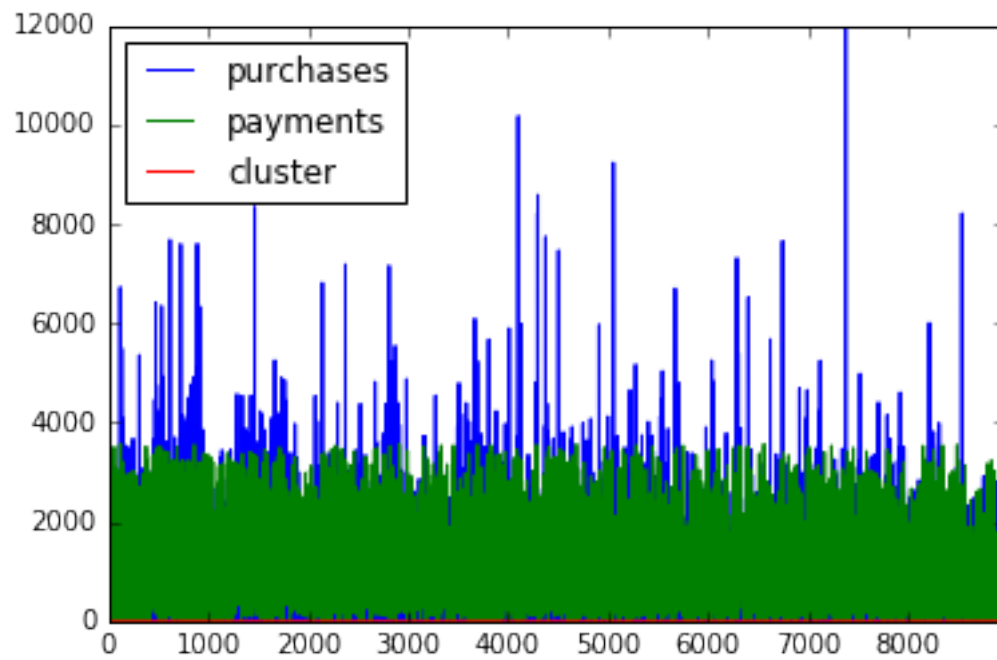
		purchases	payments	cluster
cluster				
0	count	977.000000	977.000000	977
	mean	2893.507165	6159.668584	0
	std	3248.509337	2505.563838	0

	min	0.000000	3569.182969	0
	25%	147.120000	4201.694496	0
	50%	1771.600000	5359.784024	0
	75%	4587.210000	7443.805810	0
	max	22500.000000	14481.465440	0
1	count	7899.000000	7899.000000	7899
	mean	680.066701	980.439425	1
	std	965.459756	829.030774	0
	min	0.000000	0.000000	1
	25%	35.000000	341.445479	1
	50%	324.000000	708.699641	1
	75%	924.645000	1411.360861	1
	max	11994.710000	3567.009988	1
2	count	74.000000	74.000000	74
	mean	10538.917432	23637.165609	2
	std	12900.044316	7698.746165	0
	min	0.000000	15043.665080	2
	25%	30.000000	17827.712822	2
	50%	3976.770000	20907.273725	2
	75%	21538.785000	27693.619535	2
	max	49039.570000	50721.483360	2

In [20]: grouped_cluster_data.plot(subplots=True,)

Out[20]: cluster
 0 Axes(0.125,0.125;0.775x0.775)
 1 Axes(0.125,0.125;0.775x0.775)
 2 Axes(0.125,0.125;0.775x0.775)
 dtype: object





1.6 In - class assignment

1. Provide three(3) plots of the data to assist in describing the initial data set
2. Plot the differences between the groups above using at least two (2) charts.
3. Repeat the clustering activity on different columns in an attempt to provide additional marketing insight. If the results are not insightful state why